

Final Project Individual Report: Machine Learning II
Submitted to: Dr. Amir Jafari
By Manohar Farmer

Project Title

Image Classification on Online Stores Fashion Dataset

Introduction. An overview of the project and an outline of the shared work.

The aim of this project is to build a Convolution Neural Network to Identify different types of online clothing items. The model will be evaluated with the following metrics: accuracy, F1 score, precision and recall, as well as AUC and ROC curves. To achieve this we will be tuning hyperparameters such as the number of epochs, number of layers, kernel size, etc. to improve our model. We did most of this project together as a group since we were able to meet. This included the some of the report. And the data processing as well as decisions about which frameworks, and architecture to use for our model.

Description of your individual work.

I worked on the report heavily, as well as deciding on the structure of our models. I started this process by using the MNIST code found in the github, but using the CIFAR dataset. When I was able to get the original two-layer code running well, I then began experimenting with adding layers and dropout nodes. In this experimentation dropout nodes did not improve the model, however, adding layers did. I also used this code to observe the behavior that occurred when adding a bottleneck feature and found that for that dataset, it did perform better. This result, however, was not consistent when we made these attempts on our clothing dataset.

Portion of work.

As stated above, I helped build our original algorithms to practice running the data in pytorch. I also wrote a script that we were originally going to use for determining how to resize our images, but we ended up using an arbitrary number for this value instead. I am responsible for the wording of most of the report and conducted the necessary steps on github to upload our project.

Results.

With Channel of 3 and the same random sample overtime with set.seed (1122), our results are summarized as below.

Epoch	Iteration	Batchsize	Layers	Accuracy	Time	Image Size	Optimizer
1	4	100	2	65%	--	28x28	Adam

1	4	100	3	56%	104.281539 sec	224x22 4	Adam
2	7	64		65%	11121.98764 6 sec	224x22 4	Adam
4	7	64	5 (bottleneck) 4 (maxpool)	70%	422.83998 sec	224x22 4	Adam
5	7	64	5 (bottleneck) 4 (maxpool)	72%	717.42504 sec	224x22 4	Adam
5	7	64	5 (straight) 4 (maxpool)	76%	795.74 sec	224x22 4	Adam
5	7	64	5 (bottleneck) 4 (maxpool)	80%	473.503 sec	224x22 4	SGD
5	7	64	5 (bottleneck) 4 (maxpool)	79%	481.23 sec	224x22 4	Adagrad
10	7	64	5 (bottleneck) 4 (maxpool)	79%	1542.556 sec	224x22 4	SGD
10	7	64	5 (bottleneck) 4 (maxpool)	81%	972.23 sec	224x22 4	Adagrad
5	7	64	5 (bottleneck) 4 (maxpool)	54%	1025.029 sec	224x22 4	Rprop
10	7	64	7 (straight) 4 (maxpool)	71%	1727.82sec	224x22 4	Adam
10	7	64	7(bottleneck) 4(maxpool)	69%	1511.426 sec	224x22 4	Adam

With channel of 3 and set.seed to 1122, we tuned different hyperparameters to see changes that can cause to the results.

From the results we can interpret that the accuracy rate goes up when the layers are “straight” instead of bottleneck. Also when comparing the time it takes, we find that the time it takes increases when the layers are straight. There seems to be about 4% differences between straight and bottleneck approach. Therefore, for this dataset, the straight approach performs better than the bottleneck approach and we find that straight approach isn’t preferable in this image classification. The models that have the “straight” specification were run with identical layers with the exception of the max-pooling function which were only applied to the the four outermost layers to avoid feature map size issues.


Changing kernel-size resulted in a lower accuracy of the model for our 7 layer model. The models that have the “straight” specification were run with identical layers with the exception of the max-pooling function which were only applied to the the four outermost layers to avoid feature map size issues. Changing the optimizer to SGD changes the result of the performance and increase the performance of our model. The model accuracy rate increases from 72% to 80%, which is a significant 8% increase.

The total time to run the model decreases with the change in the optimizer as well. The total time with the Adam optimizer was 795.74 sec and the total time with the SGD optimizer was 473.50 while keeping all the other parameters the same, which shows that SGD optimizer not only increases the performance of the model accuracy but also saves the time it took to run the model.

SGD, Stochastic gradient descent optimizer is called stochastic because samples are selected randomly instead of a single group. While it performs better than Adam, increasing the layer using the SGD doesn’t improve the model performance, and the total time took longer.


Confusion Matrix

The bottom two graphs are confusion matrix obtained comparing the bottleneck and straight approach with 5 layers convolutional neural network with Adam optimizer. From the diagonal line, we can observed that the straight approach with 5 layers has better accuracy. (5 layers left-bottleneck, right-straight)




When comparing the confusion matrix models using of SGD optimizer and Adam optimizer with straight technique, we find that SGD optimizer has accuracy scores that are higher than the accuracy scores of Adam optimizer.

(straight SGD 5 layers, and straight Adam 7 layers)



Accuracy vs Epoch

From the bottleneck approach with 5 layers, we also observed that the accuracy rates increase with the increase in the number of epoch from 0.45 to 0.65 in this case.




The ROC curve

Receiver Operating Characteristic (ROC) curve is defined as “the true positive rate (Sensitivity) is plotted function of the false positive rate (100-Specificity) for different cut-off


points”. From the two graphs below, we can see that the right graph has the curve which is closer to the upper corner and we can interpret that it has more accuracy. On the other hand, we can also interpret using the area under the curve (AUC) to determine and compare the accuracy of this bottleneck and straight approach. The areas under the graphs are specified in the box located at the upper corner for each item, and we find that the footwears has the highest accuracy in our image classification when comparing the 5 layers with bottleneck technique and straight technique.

(left-bottleneck, right-straight)



The two graphs below shows the SGD and Adam ROC graphs and the graphs show that SGD performs better than Adam optimizer.

(left-SGD, right-Adam)



Accuracy rates on each class (SGD best model)

Test Accuracy of the model on the test images: 80 %

Accuracy of bags : 85 %
 Accuracy of dresses : 83 %
 Accuracy of footwear : 90 %
 Accuracy of outerwear : 74 %
 Accuracy of skirts : 88 %
 Accuracy of tops : 62 %

Classification Report (SGD best model)

	precision	recall	f1-score	support
bags	0.84	0.85	0.85	4137
dresses	0.81	0.79	0.80	3770
footwear	0.94	0.89	0.92	4504
outerwear	0.69	0.77	0.73	3777
skirts	0.75	0.86	0.80	3515
tops	0.79	0.65	0.71	3820
avg / total	0.81	0.81	0.81	23523

{0: 0.9078810302987728, 1: 0.8789473337898108, 2: 0.9406002873447017, 3: 0.8529724442436417, 4: 0.9042918181233909, 5: 0.8077317271428355}

Remaining Results:

5 Layers bottle neck 4 max pool, 5 epochs, 64 batchsize

images shape on batch size = torch.Size([64, 3, 224, 224])

labels shape on batch size = torch.Size([64])

Epoch [1/5], Iter [100/734] Loss: 2.1190

Epoch [1/5], Iter [200/734] Loss: 1.6336

Epoch [1/5], Iter [300/734] Loss: 1.7083

Epoch [1/5], Iter [400/734] Loss: 1.4623

Epoch [1/5], Iter [500/734] Loss: 1.0955

Epoch [1/5], Iter [600/734] Loss: 1.1057

Epoch [1/5], Iter [700/734] Loss: 1.0495

Epoch [2/5], Iter [100/734] Loss: 1.0009

Epoch [2/5], Iter [200/734] Loss: 0.5802

Epoch [2/5], Iter [300/734] Loss: 0.8031

Epoch [2/5], Iter [400/734] Loss: 0.8572

Epoch [2/5], Iter [500/734] Loss: 0.8639

Epoch [2/5], Iter [600/734] Loss: 0.6727

Epoch [2/5], Iter [700/734] Loss: 1.1910

Epoch [3/5], Iter [100/734] Loss: 0.6867

Epoch [3/5], Iter [200/734] Loss: 0.8000

Epoch [3/5], Iter [300/734] Loss: 0.6526

Epoch [3/5], Iter [400/734] Loss: 0.9673

Epoch [3/5], Iter [500/734] Loss: 0.7886

Epoch [3/5], Iter [600/734] Loss: 0.7936

Epoch [3/5], Iter [700/734] Loss: 0.8199

Epoch [4/5], Iter [100/734] Loss: 0.9262

Epoch [4/5], Iter [200/734] Loss: 0.5819

Epoch [4/5], Iter [300/734] Loss: 0.5535

Epoch [4/5], Iter [400/734] Loss: 0.9992

Epoch [4/5], Iter [500/734] Loss: 0.7444

Epoch [4/5], Iter [600/734] Loss: 0.4865

Epoch [4/5], Iter [700/734] Loss: 0.5948

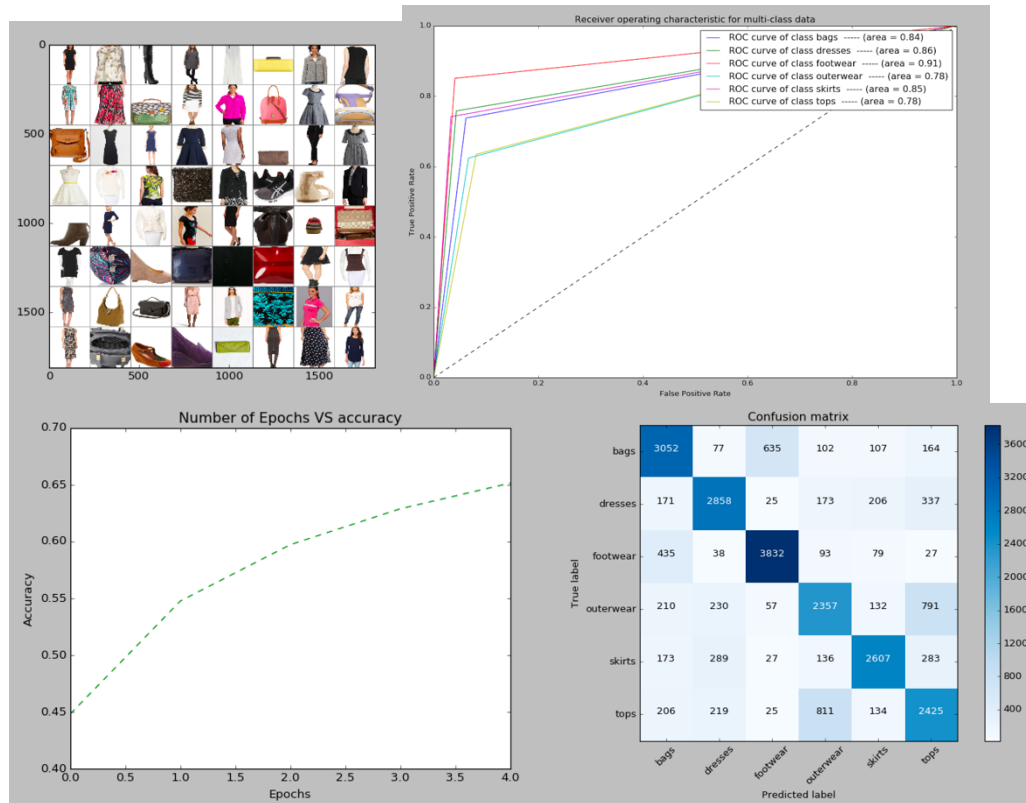
Epoch [5/5], Iter [100/734] Loss: 0.6216

Epoch [5/5], Iter [200/734] Loss: 0.7800

Epoch [5/5], Iter [300/734] Loss: 0.6399
 Epoch [5/5], Iter [400/734] Loss: 0.4969
 Epoch [5/5], Iter [500/734] Loss: 0.8024
 Epoch [5/5], Iter [600/734] Loss: 0.5087
 Epoch [5/5], Iter [700/734] Loss: 0.5625
 Training set Accuracy of the model on the Train images: 65 %
 717.424504 seconds
 Test Accuracy of the model on the test images: 72 %
 Accuracy of bags : 74 %
 Accuracy of dresses : 73 %
 Accuracy of footwear : 86 %
 Accuracy of outerwear : 65 %
 Accuracy of skirts : 75 %
 Accuracy of tops : 66 %

	precision	recall	f1-score	support
bags	0.72	0.74	0.73	4137
dresses	0.77	0.76	0.76	3770
footwear	0.83	0.85	0.84	4504
outerwear	0.64	0.62	0.63	3777
skirts	0.80	0.74	0.77	3515
tops	0.60	0.63	0.62	3820
avg / total	0.73	0.73	0.73	23523

{0: 0.8380451170738629, 1: 0.857453435874697, 2: 0.9051830192804298, 3: 0.7787222386742125, 4: 0.8543958376819968, 5: 0.776754669406126}



5 Layers straight 4 max pool, 5 epochs, 64 batchsize

images shape on batch size = torch.Size([64, 3, 224, 224])

labels shape on batch size = torch.Size([64])

bags footwear outerwear tops outerwear bags

Epoch [1/5], Iter [100/734] Loss: 2.3676

Epoch [1/5], Iter [200/734] Loss: 1.6400

Epoch [1/5], Iter [300/734] Loss: 1.1064

Epoch [1/5], Iter [400/734] Loss: 1.2984

Epoch [1/5], Iter [500/734] Loss: 1.1671

Epoch [1/5], Iter [600/734] Loss: 1.0433

Epoch [1/5], Iter [700/734] Loss: 0.8376

Epoch [2/5], Iter [100/734] Loss: 1.0417

Epoch [2/5], Iter [200/734] Loss: 0.7722

Epoch [2/5], Iter [300/734] Loss: 0.8563

Epoch [2/5], Iter [400/734] Loss: 1.0028

Epoch [2/5], Iter [500/734] Loss: 0.6126

Epoch [2/5], Iter [600/734] Loss: 0.9350

Epoch [2/5], Iter [700/734] Loss: 0.6436

Epoch [3/5], Iter [100/734] Loss: 0.7296

Epoch [3/5], Iter [200/734] Loss: 0.4224

Epoch [3/5], Iter [300/734] Loss: 0.5868

Epoch [3/5], Iter [400/734] Loss: 0.5187


Epoch [3/5], Iter [500/734] Loss: 0.7172

Epoch [3/5], Iter [600/734] Loss: 0.8649

Epoch [3/5], Iter [700/734] Loss: 0.7659
 Epoch [4/5], Iter [100/734] Loss: 0.6213
 Epoch [4/5], Iter [200/734] Loss: 0.9780
 Epoch [4/5], Iter [300/734] Loss: 0.8408
 Epoch [4/5], Iter [400/734] Loss: 0.5910
 Epoch [4/5], Iter [500/734] Loss: 0.5645
 Epoch [4/5], Iter [600/734] Loss: 0.5254
 Epoch [4/5], Iter [700/734] Loss: 0.6211
 Epoch [5/5], Iter [100/734] Loss: 0.6844
 Epoch [5/5], Iter [200/734] Loss: 0.7029
 Epoch [5/5], Iter [300/734] Loss: 0.5809
 Epoch [5/5], Iter [400/734] Loss: 0.8297
 Epoch [5/5], Iter [500/734] Loss: 0.5689
 Epoch [5/5], Iter [600/734] Loss: 0.3840
 Epoch [5/5], Iter [700/734] Loss: 0.5465
 Training set Accuracy of the model on the Train images: 67 %
 795.74182 seconds
 Test Accuracy of the model on the test images: 76 %
 Accuracy of bags : 85 %
 Accuracy of dresses : 78 %
 Accuracy of footwear : 88 %
 Accuracy of outerwear : 59 %
 Accuracy of skirts : 80 %
 Accuracy of tops : 75 %

	precision	recall	f1-score	support
bags	0.75	0.82	0.78	4137
dresses	0.75	0.77	0.76	3770
footwear	0.87	0.88	0.88	4504
outerwear	0.72	0.59	0.65	3777
skirts	0.81	0.76	0.78	3515
tops	0.66	0.71	0.68	3820
avg / total	0.76	0.76	0.76	23523

{0: 0.8807220875960889, 1: 0.8596379007533489, 2: 0.9240341433830263, 3:
 0.7740909430626783, 4: 0.8642493784847369, 5: 0.8199951744133365}
 GroundTruth: bags bags bags bags
 Predicted: bags footwear bags outerwear



SGD

images shape on batch size = torch.Size([64, 3, 224, 224])

labels shape on batch size = torch.Size([64])

Epoch [1/5], Iter [100/734] Loss: 1.1525

Epoch [1/5], Iter [200/734] Loss: 0.8715

Epoch [1/5], Iter [300/734] Loss: 0.8618

Epoch [1/5], Iter [400/734] Loss: 0.6990

Epoch [1/5], Iter [500/734] Loss: 0.7519

Epoch [1/5], Iter [600/734] Loss: 0.8176

Epoch [1/5], Iter [700/734] Loss: 0.8236

Epoch [2/5], Iter [100/734] Loss: 0.5339

Epoch [2/5], Iter [200/734] Loss: 0.5881

Epoch [2/5], Iter [300/734] Loss: 0.7327

Epoch [2/5], Iter [400/734] Loss: 0.5550

Epoch [2/5], Iter [500/734] Loss: 0.6153

Epoch [2/5], Iter [600/734] Loss: 0.5799

Epoch [2/5], Iter [700/734] Loss: 0.5511

Epoch [3/5], Iter [100/734] Loss: 0.4423

Epoch [3/5], Iter [200/734] Loss: 0.6513

Epoch [3/5], Iter [300/734] Loss: 0.3372

Epoch [3/5], Iter [400/734] Loss: 0.5801

Epoch [3/5], Iter [500/734] Loss: 0.5338

Epoch [3/5], Iter [600/734] Loss: 0.7010

Epoch [3/5], Iter [700/734] Loss: 0.5461

Epoch [4/5], Iter [100/734] Loss: 0.4753

Epoch [4/5], Iter [200/734] Loss: 0.3725

Epoch [4/5], Iter [300/734] Loss: 0.3751

Epoch [4/5], Iter [400/734] Loss: 0.6334

Epoch [4/5], Iter [500/734] Loss: 0.4577

Epoch [4/5], Iter [600/734] Loss: 0.6039

Epoch [4/5], Iter [700/734] Loss: 0.3366

Epoch [5/5], Iter [100/734] Loss: 0.3060

Epoch [5/5], Iter [200/734] Loss: 0.4821

Epoch [5/5], Iter [300/734] Loss: 0.8020

Epoch [5/5], Iter [400/734] Loss: 0.4375
Epoch [5/5], Iter [500/734] Loss: 0.5330
Epoch [5/5], Iter [600/734] Loss: 0.5881
Epoch [5/5], Iter [700/734] Loss: 0.5916
473.502485 seconds
Test Accuracy of the model on the test images: 80 %

Accuracy of bags : 85 %
Accuracy of dresses : 83 %
Accuracy of footwear : 90 %
Accuracy of outerwear : 74 %
Accuracy of skirts : 88 %
Accuracy of tops : 62 %

	precision	recall	f1-score	support
bags	0.84	0.85	0.85	4137
dresses	0.81	0.79	0.80	3770
footwear	0.94	0.89	0.92	4504
outerwear	0.69	0.77	0.73	3777
skirts	0.75	0.86	0.80	3515
tops	0.79	0.65	0.71	3820
avg / total	0.81	0.81	0.81	23523

{0: 0.9078810302987728, 1: 0.8789473337898108, 2: 0.9406002873447017, 3: 0.8529724442436417, 4: 0.9042918181233909, 5: 0.8077317271428355}

7Layers straight 4maxpool 10epochs 64batchsize

cuda:0

```
<class 'torch.utils.data.dataloader._DataLoaderIter'>  
images shape on batch size = torch.Size([64, 3, 224, 224])  
labels shape on batch size = torch.Size([64])  
Epoch [1/10], Iter [100/734] Loss: 1.7384  
Epoch [1/10], Iter [200/734] Loss: 1.5978  
Epoch [1/10], Iter [300/734] Loss: 1.6892  
Epoch [1/10], Iter [400/734] Loss: 1.1238  
Epoch [1/10], Iter [500/734] Loss: 1.1008  
Epoch [1/10], Iter [600/734] Loss: 1.7555  
Epoch [1/10], Iter [700/734] Loss: 1.2471  
Epoch [2/10], Iter [100/734] Loss: 1.2666  
Epoch [2/10], Iter [200/734] Loss: 1.0766  
Epoch [2/10], Iter [300/734] Loss: 1.1622  
Epoch [2/10], Iter [400/734] Loss: 0.8324  
Epoch [2/10], Iter [500/734] Loss: 1.0236  
Epoch [2/10], Iter [600/734] Loss: 1.4816  
Epoch [2/10], Iter [700/734] Loss: 1.0749  
Epoch [3/10], Iter [100/734] Loss: 1.0550  
Epoch [3/10], Iter [200/734] Loss: 0.8160
```

Epoch [3/10], Iter [300/734] Loss: 0.9476
Epoch [3/10], Iter [400/734] Loss: 0.8308
Epoch [3/10], Iter [500/734] Loss: 0.8286
Epoch [3/10], Iter [600/734] Loss: 1.0995
Epoch [3/10], Iter [700/734] Loss: 0.8181
Epoch [4/10], Iter [100/734] Loss: 0.9433
Epoch [4/10], Iter [200/734] Loss: 0.9808
Epoch [4/10], Iter [300/734] Loss: 0.9997
Epoch [4/10], Iter [400/734] Loss: 0.8923
Epoch [4/10], Iter [500/734] Loss: 0.9113
Epoch [4/10], Iter [600/734] Loss: 0.9280
Epoch [4/10], Iter [700/734] Loss: 0.9028
Epoch [5/10], Iter [100/734] Loss: 0.7001
Epoch [5/10], Iter [200/734] Loss: 0.9468
Epoch [5/10], Iter [300/734] Loss: 0.9191
Epoch [5/10], Iter [400/734] Loss: 0.8411
Epoch [5/10], Iter [500/734] Loss: 0.9765
Epoch [5/10], Iter [600/734] Loss: 0.9019
Epoch [5/10], Iter [700/734] Loss: 0.8397
Epoch [6/10], Iter [100/734] Loss: 0.7925
Epoch [6/10], Iter [200/734] Loss: 1.1802
Epoch [6/10], Iter [300/734] Loss: 0.9014
Epoch [6/10], Iter [400/734] Loss: 0.6037
Epoch [6/10], Iter [500/734] Loss: 0.9972
Epoch [6/10], Iter [600/734] Loss: 0.9715
Epoch [6/10], Iter [700/734] Loss: 0.8063
Epoch [7/10], Iter [100/734] Loss: 0.6520
Epoch [7/10], Iter [200/734] Loss: 1.0989
Epoch [7/10], Iter [300/734] Loss: 0.8753
Epoch [7/10], Iter [400/734] Loss: 0.7345
Epoch [7/10], Iter [500/734] Loss: 0.6590
Epoch [7/10], Iter [600/734] Loss: 0.7322
Epoch [7/10], Iter [700/734] Loss: 0.7358
Epoch [8/10], Iter [100/734] Loss: 0.8404
Epoch [8/10], Iter [200/734] Loss: 0.5090
Epoch [8/10], Iter [300/734] Loss: 0.8169
Epoch [8/10], Iter [400/734] Loss: 0.6299
Epoch [8/10], Iter [500/734] Loss: 0.6758
Epoch [8/10], Iter [600/734] Loss: 0.7371
Epoch [8/10], Iter [700/734] Loss: 0.5271
Epoch [9/10], Iter [100/734] Loss: 0.5350
Epoch [9/10], Iter [200/734] Loss: 0.8278
Epoch [9/10], Iter [300/734] Loss: 1.0682
Epoch [9/10], Iter [400/734] Loss: 1.0672
Epoch [9/10], Iter [500/734] Loss: 0.7146
Epoch [9/10], Iter [600/734] Loss: 0.6417
Epoch [9/10], Iter [700/734] Loss: 0.8504
Epoch [10/10], Iter [100/734] Loss: 0.6552
Epoch [10/10], Iter [200/734] Loss: 0.6977
Epoch [10/10], Iter [300/734] Loss: 0.7299
Epoch [10/10], Iter [400/734] Loss: 0.7889

Epoch [10/10], Iter [500/734] Loss: 0.8208
 Epoch [10/10], Iter [600/734] Loss: 0.5978
 Epoch [10/10], Iter [700/734] Loss: 0.6368
 Training set Accuracy of the model on the Train images: 64 %
 1727.824153 seconds
 Test Accuracy of the model on the test images: 71 %
 Accuracy of bags : 81 %
 Accuracy of dresses : 74 %
 Accuracy of footwear : 76 %
 Accuracy of outerwear : 63 %
 Accuracy of skirts : 81 %
 Accuracy of tops : 52 %

	precision	recall	f1-score	support
bags	0.72	0.80	0.76	4137
dresses	0.72	0.78	0.75	3770
footwear	0.91	0.75	0.82	4504
outerwear	0.58	0.66	0.62	3777
skirts	0.71	0.78	0.74	3515
tops	0.69	0.55	0.61	3820
avg / total	0.73	0.72	0.72	23523

{0: 0.8663738370587628, 1: 0.8597762875491094, 2: 0.8648170155076297, 3:
 0.7844414372297841, 4: 0.8623133747923306, 5: 0.7515264172968583}

7Layers straight 4maxpool 10epochs 64batchsize 3kernelsize

images shape on batch size = torch.Size([64, 3, 224, 224])

labels shape on batch size = torch.Size([64])

Epoch [1/10], Iter [100/734] Loss: 1.5868
 Epoch [1/10], Iter [200/734] Loss: 1.5498
 Epoch [1/10], Iter [300/734] Loss: 1.3855
 Epoch [1/10], Iter [400/734] Loss: 1.4096
 Epoch [1/10], Iter [500/734] Loss: 1.4473
 Epoch [1/10], Iter [600/734] Loss: 1.3095
 Epoch [1/10], Iter [700/734] Loss: 1.2992
 Epoch [2/10], Iter [100/734] Loss: 1.1442
 Epoch [2/10], Iter [200/734] Loss: 1.2236
 Epoch [2/10], Iter [300/734] Loss: 0.9719
 Epoch [2/10], Iter [400/734] Loss: 0.8542
 Epoch [2/10], Iter [500/734] Loss: 1.0130
 Epoch [2/10], Iter [600/734] Loss: 0.8867
 Epoch [2/10], Iter [700/734] Loss: 0.9214
 Epoch [3/10], Iter [100/734] Loss: 0.8348
 Epoch [3/10], Iter [200/734] Loss: 0.7168
 Epoch [3/10], Iter [300/734] Loss: 0.9214
 Epoch [3/10], Iter [400/734] Loss: 0.6325
 Epoch [3/10], Iter [500/734] Loss: 0.8184

Epoch [3/10], Iter [600/734] Loss: 0.9593
Epoch [3/10], Iter [700/734] Loss: 0.7305
Epoch [4/10], Iter [100/734] Loss: 0.7207
Epoch [4/10], Iter [200/734] Loss: 0.8923
Epoch [4/10], Iter [300/734] Loss: 0.9438
Epoch [4/10], Iter [400/734] Loss: 0.6959
Epoch [4/10], Iter [500/734] Loss: 0.7747
Epoch [4/10], Iter [600/734] Loss: 0.7402
Epoch [4/10], Iter [700/734] Loss: 0.7599
Epoch [5/10], Iter [100/734] Loss: 0.7755
Epoch [5/10], Iter [200/734] Loss: 0.7499
Epoch [5/10], Iter [300/734] Loss: 0.7914
Epoch [5/10], Iter [400/734] Loss: 0.9518
Epoch [5/10], Iter [500/734] Loss: 0.6999
Epoch [5/10], Iter [600/734] Loss: 0.7043
Epoch [5/10], Iter [700/734] Loss: 0.7256
Epoch [6/10], Iter [100/734] Loss: 0.4863
Epoch [6/10], Iter [200/734] Loss: 0.8476
Epoch [6/10], Iter [300/734] Loss: 0.8546
Epoch [6/10], Iter [400/734] Loss: 0.6579
Epoch [6/10], Iter [500/734] Loss: 0.5809
Epoch [6/10], Iter [600/734] Loss: 0.5134
Epoch [6/10], Iter [700/734] Loss: 0.7616
Epoch [7/10], Iter [100/734] Loss: 0.5119
Epoch [7/10], Iter [200/734] Loss: 0.6025
Epoch [7/10], Iter [300/734] Loss: 0.6867
Epoch [7/10], Iter [400/734] Loss: 0.6755
Epoch [7/10], Iter [500/734] Loss: 0.4947
Epoch [7/10], Iter [600/734] Loss: 0.5777
Epoch [7/10], Iter [700/734] Loss: 0.7827
Epoch [8/10], Iter [100/734] Loss: 0.6306
Epoch [8/10], Iter [200/734] Loss: 0.4200
Epoch [8/10], Iter [300/734] Loss: 0.5762
Epoch [8/10], Iter [400/734] Loss: 0.5348
Epoch [8/10], Iter [500/734] Loss: 0.5426
Epoch [8/10], Iter [600/734] Loss: 0.6827
Epoch [8/10], Iter [700/734] Loss: 0.5600
Epoch [9/10], Iter [100/734] Loss: 0.7473
Epoch [9/10], Iter [200/734] Loss: 0.4525
Epoch [9/10], Iter [300/734] Loss: 0.5563
Epoch [9/10], Iter [400/734] Loss: 0.6238
Epoch [9/10], Iter [500/734] Loss: 0.6597
Epoch [9/10], Iter [600/734] Loss: 0.7225
Epoch [9/10], Iter [700/734] Loss: 0.4683
Epoch [10/10], Iter [100/734] Loss: 0.6528
Epoch [10/10], Iter [200/734] Loss: 0.8435
Epoch [10/10], Iter [300/734] Loss: 0.6604
Epoch [10/10], Iter [400/734] Loss: 0.4524
Epoch [10/10], Iter [500/734] Loss: 0.4923
Epoch [10/10], Iter [600/734] Loss: 0.6113
Epoch [10/10], Iter [700/734] Loss: 0.5534

Training set Accuracy of the model on the Train images: 67 %

1511.4266280000002 seconds

Test Accuracy of the model on the test images: 69 %

Accuracy of bags : 76 %

Accuracy of dresses : 76 %

Accuracy of footwear : 85 %

Accuracy of outerwear : 58 %

Accuracy of skirts : 84 %

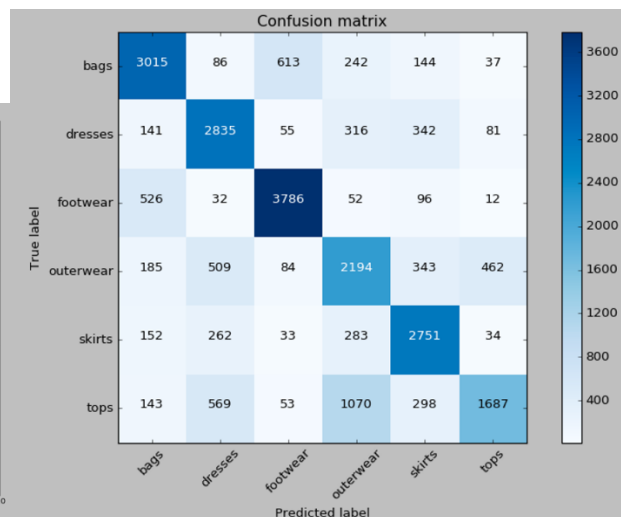
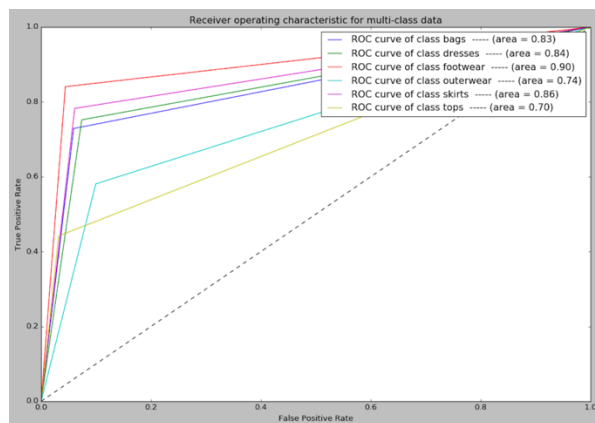
Accuracy of tops : 40 %

precision recall f1-score support

bags	0.72	0.73	0.73	4137
dresses	0.66	0.75	0.70	3770
footwear	0.82	0.84	0.83	4504
outerwear	0.53	0.58	0.55	3777
skirts	0.69	0.78	0.73	3515
tops	0.73	0.44	0.55	3820

avg / total 0.70 0.69 0.69 23523

{0: 0.8348112843856803, 1: 0.8390889084973965, 2: 0.8982624718461869, 3: 0.740735880230154, 4: 0.8607601269591736, 5: 0.7049256126249677}



7Layers bottleneck 4maxpool 10epochs 64batchsize

cuda:0

<class 'torch.utils.data.dataloader._DataLoaderIter'>

images shape on batch size = torch.Size([64, 3, 224, 224])

labels shape on batch size = torch.Size([64])

tops tops footwear skirts bags tops footwear outerwear bags skirts footwear dresses

footwear bags outerwear dresses skirts outerwear skirts footwear outerwear tops footwear

footwear skirts outerwear dresses outerwear tops skirts footwear skirts footwear bags

outerwear outerwear tops skirts skirts outerwear bags dresses footwear tops skirts bags

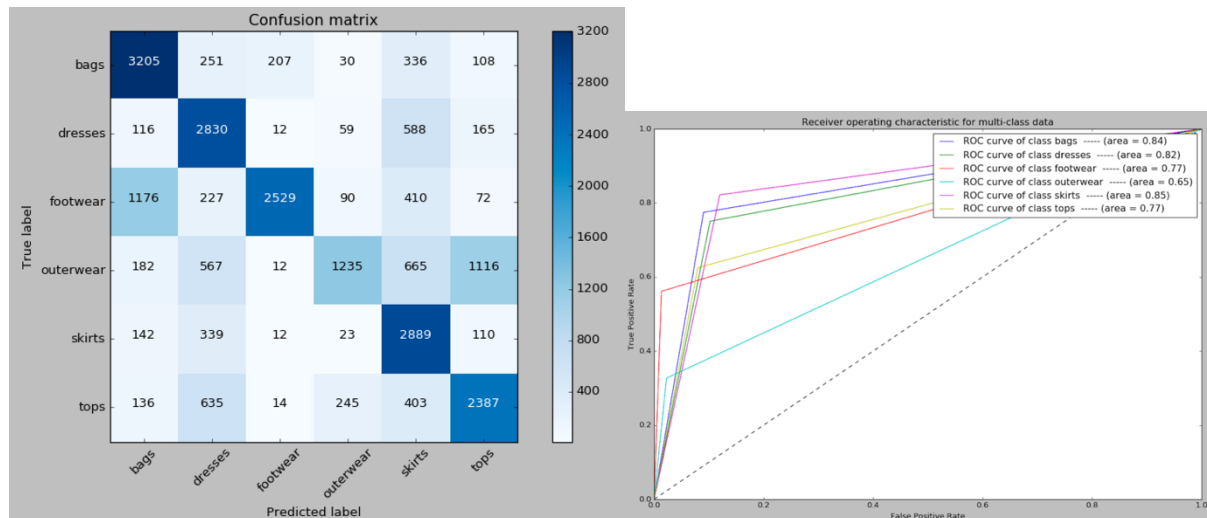
dresses bags bags bags footwear outerwear skirts outerwear outerwear tops skirts bags
bags footwear outerwear tops outerwear bags

Epoch [1/10], Iter [100/734] Loss: 1.4676
Epoch [1/10], Iter [200/734] Loss: 1.3276
Epoch [1/10], Iter [300/734] Loss: 1.2785
Epoch [1/10], Iter [400/734] Loss: 1.1935
Epoch [1/10], Iter [500/734] Loss: 1.3077
Epoch [1/10], Iter [600/734] Loss: 0.9357
Epoch [1/10], Iter [700/734] Loss: 1.0591
Epoch [2/10], Iter [100/734] Loss: 0.8426
Epoch [2/10], Iter [200/734] Loss: 0.8500
Epoch [2/10], Iter [300/734] Loss: 0.9319
Epoch [2/10], Iter [400/734] Loss: 1.2197
Epoch [2/10], Iter [500/734] Loss: 0.7788
Epoch [2/10], Iter [600/734] Loss: 0.8853
Epoch [2/10], Iter [700/734] Loss: 0.8381
Epoch [3/10], Iter [100/734] Loss: 0.7493
Epoch [3/10], Iter [200/734] Loss: 0.6120
Epoch [3/10], Iter [300/734] Loss: 0.9029
Epoch [3/10], Iter [400/734] Loss: 0.5365
Epoch [3/10], Iter [500/734] Loss: 0.8527
Epoch [3/10], Iter [600/734] Loss: 0.6645
Epoch [3/10], Iter [700/734] Loss: 0.6812
Epoch [4/10], Iter [100/734] Loss: 0.8217
Epoch [4/10], Iter [200/734] Loss: 0.7340
Epoch [4/10], Iter [300/734] Loss: 0.7476
Epoch [4/10], Iter [400/734] Loss: 0.6601
Epoch [4/10], Iter [500/734] Loss: 0.7787
Epoch [4/10], Iter [600/734] Loss: 0.7153
Epoch [4/10], Iter [700/734] Loss: 0.4962
Epoch [5/10], Iter [100/734] Loss: 0.7229
Epoch [5/10], Iter [200/734] Loss: 0.7583
Epoch [5/10], Iter [300/734] Loss: 0.6273
Epoch [5/10], Iter [400/734] Loss: 0.7988
Epoch [5/10], Iter [500/734] Loss: 0.5132
Epoch [5/10], Iter [600/734] Loss: 0.6681
Epoch [5/10], Iter [700/734] Loss: 0.9181
Epoch [6/10], Iter [100/734] Loss: 0.7092
Epoch [6/10], Iter [200/734] Loss: 0.8273
Epoch [6/10], Iter [300/734] Loss: 0.6527
Epoch [6/10], Iter [400/734] Loss: 0.6090
Epoch [6/10], Iter [500/734] Loss: 0.6638
Epoch [6/10], Iter [600/734] Loss: 0.6300
Epoch [6/10], Iter [700/734] Loss: 0.7598
Epoch [7/10], Iter [100/734] Loss: 0.4893
Epoch [7/10], Iter [200/734] Loss: 0.5426
Epoch [7/10], Iter [300/734] Loss: 0.6837
Epoch [7/10], Iter [400/734] Loss: 0.7651
Epoch [7/10], Iter [500/734] Loss: 0.4926
Epoch [7/10], Iter [600/734] Loss: 0.5789
Epoch [7/10], Iter [700/734] Loss: 0.8050

Epoch [8/10], Iter [100/734] Loss: 0.7930
 Epoch [8/10], Iter [200/734] Loss: 0.5370
 Epoch [8/10], Iter [300/734] Loss: 0.6571
 Epoch [8/10], Iter [400/734] Loss: 0.7404
 Epoch [8/10], Iter [500/734] Loss: 0.6290
 Epoch [8/10], Iter [600/734] Loss: 0.5522
 Epoch [8/10], Iter [700/734] Loss: 0.6183
 Epoch [9/10], Iter [100/734] Loss: 0.6639
 Epoch [9/10], Iter [200/734] Loss: 0.6806
 Epoch [9/10], Iter [300/734] Loss: 0.5304
 Epoch [9/10], Iter [400/734] Loss: 0.8879
 Epoch [9/10], Iter [500/734] Loss: 0.5506
 Epoch [9/10], Iter [600/734] Loss: 0.8253
 Epoch [9/10], Iter [700/734] Loss: 0.6097
 Epoch [10/10], Iter [100/734] Loss: 0.4011
 Epoch [10/10], Iter [200/734] Loss: 0.5939
 Epoch [10/10], Iter [300/734] Loss: 0.6401
 Epoch [10/10], Iter [400/734] Loss: 0.5383
 Epoch [10/10], Iter [500/734] Loss: 0.6387
 Epoch [10/10], Iter [600/734] Loss: 0.6728
 Epoch [10/10], Iter [700/734] Loss: 0.9676
 Training set Accuracy of the model on the Train images: 70 %
 1654.217491 seconds
 Test Accuracy of the model on the test images: 64 %
 Accuracy of bags : 80 %
 Accuracy of dresses : 77 %
 Accuracy of footwear : 60 %
 Accuracy of outerwear : 27 %
 Accuracy of skirts : 85 %
 Accuracy of tops : 60 %

	precision	recall	f1-score	support
bags	0.65	0.77	0.70	4137
dresses	0.58	0.75	0.66	3770
footwear	0.91	0.56	0.69	4504
outerwear	0.73	0.33	0.45	3777
skirts	0.55	0.82	0.66	3515
tops	0.60	0.62	0.61	3820
avg / total	0.68	0.64	0.63	23523

{0: 0.8421707403509646, 1: 0.8242254038972827, 2: 0.7739940425564901, 3:
 0.652170793863641, 4: 0.850927068717321, 5: 0.7725675296477296}
 ('GroundTruth: ', 'bags bags bags bags')
 ('Predicted: ', 'skirts skirts bags skirts')



5Layers straight 4maxpool 10epochs 64batchsize SGD

images shape on batch size = torch.Size([64, 3, 224, 224])

labels shape on batch size = torch.Size([64])

tops tops footwear skirts bags tops footwear outerwear bags skirts footwear dresses
 footwear bags outerwear dresses skirts outerwear skirts footwear outerwear tops footwear
 footwear skirts outerwear dresses outerwear tops skirts footwear skirts footwear bags
 outerwear outerwear tops skirts skirts outerwear bags dresses footwear tops skirts bags
 dresses bags bags bags footwear outerwear skirts outerwear outerwear tops skirts bags
 bags footwear outerwear tops outerwear bags

Epoch [1/10], Iter [100/734] Loss: 1.2654

Epoch [1/10], Iter [200/734] Loss: 0.9062

Epoch [1/10], Iter [300/734] Loss: 0.9264

Epoch [1/10], Iter [400/734] Loss: 1.0646

Epoch [1/10], Iter [500/734] Loss: 0.7315

Epoch [1/10], Iter [600/734] Loss: 0.7811

Epoch [1/10], Iter [700/734] Loss: 0.6110

Epoch [2/10], Iter [100/734] Loss: 0.6881

Epoch [2/10], Iter [200/734] Loss: 0.5268

Epoch [2/10], Iter [300/734] Loss: 0.6371

Epoch [2/10], Iter [400/734] Loss: 0.8786

Epoch [2/10], Iter [500/734] Loss: 0.4430

Epoch [2/10], Iter [600/734] Loss: 0.7251

Epoch [2/10], Iter [700/734] Loss: 0.4489

Epoch [3/10], Iter [100/734] Loss: 0.5221

Epoch [3/10], Iter [200/734] Loss: 0.4119

Epoch [3/10], Iter [300/734] Loss: 0.4001

Epoch [3/10], Iter [400/734] Loss: 0.5887

Epoch [3/10], Iter [500/734] Loss: 0.4385

Epoch [3/10], Iter [600/734] Loss: 0.7185

Epoch [3/10], Iter [700/734] Loss: 0.6209

Epoch [4/10], Iter [100/734] Loss: 0.4697


Epoch [4/10], Iter [200/734] Loss: 0.6585
Epoch [4/10], Iter [300/734] Loss: 0.6378
Epoch [4/10], Iter [400/734] Loss: 0.4570
Epoch [4/10], Iter [500/734] Loss: 0.5559
Epoch [4/10], Iter [600/734] Loss: 0.3208
Epoch [4/10], Iter [700/734] Loss: 0.4059
Epoch [5/10], Iter [100/734] Loss: 0.7248
Epoch [5/10], Iter [200/734] Loss: 0.5890
Epoch [5/10], Iter [300/734] Loss: 0.6253
Epoch [5/10], Iter [400/734] Loss: 0.6055
Epoch [5/10], Iter [500/734] Loss: 0.6216
Epoch [5/10], Iter [600/734] Loss: 0.2936
Epoch [5/10], Iter [700/734] Loss: 0.4689
Epoch [6/10], Iter [100/734] Loss: 0.3743
Epoch [6/10], Iter [200/734] Loss: 0.4792
Epoch [6/10], Iter [300/734] Loss: 0.3437
Epoch [6/10], Iter [400/734] Loss: 0.4199
Epoch [6/10], Iter [500/734] Loss: 0.5919
Epoch [6/10], Iter [600/734] Loss: 0.6372
Epoch [6/10], Iter [700/734] Loss: 0.4612
Epoch [7/10], Iter [100/734] Loss: 0.5745
Epoch [7/10], Iter [200/734] Loss: 0.3906
Epoch [7/10], Iter [300/734] Loss: 0.4917
Epoch [7/10], Iter [400/734] Loss: 0.4401
Epoch [7/10], Iter [500/734] Loss: 0.2848
Epoch [7/10], Iter [600/734] Loss: 0.4011
Epoch [7/10], Iter [700/734] Loss: 0.5174
Epoch [8/10], Iter [100/734] Loss: 0.5209
Epoch [8/10], Iter [200/734] Loss: 0.3897
Epoch [8/10], Iter [300/734] Loss: 0.6272
Epoch [8/10], Iter [400/734] Loss: 0.3681
Epoch [8/10], Iter [500/734] Loss: 0.4750
Epoch [8/10], Iter [600/734] Loss: 0.3578
Epoch [8/10], Iter [700/734] Loss: 0.3371
Epoch [9/10], Iter [100/734] Loss: 0.4898
Epoch [9/10], Iter [200/734] Loss: 0.6001
Epoch [9/10], Iter [300/734] Loss: 0.4876
Epoch [9/10], Iter [400/734] Loss: 0.4811
Epoch [9/10], Iter [500/734] Loss: 0.3513
Epoch [9/10], Iter [600/734] Loss: 0.3209
Epoch [9/10], Iter [700/734] Loss: 0.3917
Epoch [10/10], Iter [100/734] Loss: 0.4042
Epoch [10/10], Iter [200/734] Loss: 0.4195
Epoch [10/10], Iter [300/734] Loss: 0.3467
Epoch [10/10], Iter [400/734] Loss: 0.5148
Epoch [10/10], Iter [500/734] Loss: 0.3263
Epoch [10/10], Iter [600/734] Loss: 0.3237
Epoch [10/10], Iter [700/734] Loss: 0.5263
Training set Accuracy of the model on the Train images: 80 %
1542.561696 seconds
Test Accuracy of the model on the test images: 79 %

Accuracy of bags : 94 %
 Accuracy of dresses : 78 %
 Accuracy of footwear : 86 %
 Accuracy of outerwear : 69 %
 Accuracy of skirts : 86 %
 Accuracy of tops : 73 %

Confusion matrix

	precision	recall	f1-score	support
bags	0.71	0.92	0.80	4137
dresses	0.85	0.78	0.81	3770
footwear	0.95	0.84	0.90	4504
outerwear	0.77	0.67	0.72	3777
skirts	0.80	0.82	0.81	3515
tops	0.73	0.73	0.73	3820
avg / total	0.81	0.80	0.80	23523

{0: 0.9213215363583702, 1: 0.8753171696982939, 2: 0.9173693056966404, 3: 0.8173811174754972, 4: 0.8938310379973188, 5: 0.8365288008603149}



images shape on batch size = torch.Size([64, 3, 224, 224])

labels shape on batch size = torch.Size([64])

Epoch [1/5], Iter [100/734] Loss: 1.6147

Epoch [1/5], Iter [200/734] Loss: 1.3458

Epoch [1/5], Iter [300/734] Loss: 1.6286

Epoch [1/5], Iter [400/734] Loss: 1.4943

Epoch [1/5], Iter [500/734] Loss: 1.6882

Epoch [1/5], Iter [600/734] Loss: 1.3587

Epoch [1/5], Iter [700/734] Loss: 1.8820

Epoch [2/5], Iter [100/734] Loss: 1.2920

Epoch [2/5], Iter [200/734] Loss: 1.2436

Epoch [2/5], Iter [300/734] Loss: 1.2560

Epoch [2/5], Iter [400/734] Loss: 1.1185

Epoch [2/5], Iter [500/734] Loss: 1.1285

Epoch [2/5], Iter [600/734] Loss: 1.5031
 Epoch [2/5], Iter [700/734] Loss: 1.1962
 Epoch [3/5], Iter [100/734] Loss: 1.2916
 Epoch [3/5], Iter [200/734] Loss: 0.9956
 Epoch [3/5], Iter [300/734] Loss: 1.3433
 Epoch [3/5], Iter [400/734] Loss: 1.0013
 Epoch [3/5], Iter [500/734] Loss: 1.1175
 Epoch [3/5], Iter [600/734] Loss: 0.9673
 Epoch [3/5], Iter [700/734] Loss: 1.3639
 Epoch [4/5], Iter [100/734] Loss: 1.6038
 Epoch [4/5], Iter [200/734] Loss: 1.1203
 Epoch [4/5], Iter [300/734] Loss: 1.0693
 Epoch [4/5], Iter [400/734] Loss: 1.1472
 Epoch [4/5], Iter [500/734] Loss: 1.5915
 Epoch [4/5], Iter [600/734] Loss: 0.9166
 Epoch [4/5], Iter [700/734] Loss: 1.0970
 Epoch [5/5], Iter [100/734] Loss: 1.3605
 Epoch [5/5], Iter [200/734] Loss: 1.0445
 Epoch [5/5], Iter [300/734] Loss: 1.2392
 Epoch [5/5], Iter [400/734] Loss: 1.1143
 Epoch [5/5], Iter [500/734] Loss: 1.4075
 Epoch [5/5], Iter [600/734] Loss: 1.1985
 Epoch [5/5], Iter [700/734] Loss: 1.4106

1025.029345 seconds

Test Accuracy of the model on the test images: 57 %

Accuracy of bags : 47 %

Accuracy of dresses : 59 %

Accuracy of footwear : 77 %

Accuracy of outerwear : 49 %

Accuracy of skirts : 67 %

Accuracy of tops : 35 %

	precision	recall	f1-score	support
bags	0.64	0.47	0.54	4137
dresses	0.55	0.61	0.58	3770
footwear	0.66	0.75	0.70	4504
outerwear	0.45	0.51	0.48	3777
skirts	0.62	0.68	0.65	3515
tops	0.52	0.41	0.46	3820
avg / total	0.58	0.57	0.57	23523

{0: 0.705917890253255, 1: 0.7575016439768542, 2: 0.8302203662468223, 3: 0.6959040859423011, 4: 0.8032823997001485, 5: 0.6680424991755846}

images shape on batch size = torch.Size([64, 3, 224, 224])

labels shape on batch size = torch.Size([64])

Epoch [1/10], Iter [100/734] Loss: 1.9546

Epoch [1/10], Iter [200/734] Loss: 1.2948

Epoch [1/10], Iter [300/734] Loss: 1.3810
Epoch [1/10], Iter [400/734] Loss: 1.0642
Epoch [1/10], Iter [500/734] Loss: 1.0350
Epoch [1/10], Iter [600/734] Loss: 0.8895
Epoch [1/10], Iter [700/734] Loss: 1.0499
Epoch [2/10], Iter [100/734] Loss: 0.6745
Epoch [2/10], Iter [200/734] Loss: 0.8023
Epoch [2/10], Iter [300/734] Loss: 0.9847
Epoch [2/10], Iter [400/734] Loss: 0.7109
Epoch [2/10], Iter [500/734] Loss: 0.7118
Epoch [2/10], Iter [600/734] Loss: 0.6310
Epoch [2/10], Iter [700/734] Loss: 0.7367
Epoch [3/10], Iter [100/734] Loss: 0.6599
Epoch [3/10], Iter [200/734] Loss: 0.7837
Epoch [3/10], Iter [300/734] Loss: 0.3868
Epoch [3/10], Iter [400/734] Loss: 0.5837
Epoch [3/10], Iter [500/734] Loss: 0.7627
Epoch [3/10], Iter [600/734] Loss: 0.7908
Epoch [3/10], Iter [700/734] Loss: 0.7715
Epoch [4/10], Iter [100/734] Loss: 0.6846
Epoch [4/10], Iter [200/734] Loss: 0.3889
Epoch [4/10], Iter [300/734] Loss: 0.3653
Epoch [4/10], Iter [400/734] Loss: 0.7146
Epoch [4/10], Iter [500/734] Loss: 0.5413
Epoch [4/10], Iter [600/734] Loss: 0.7096
Epoch [4/10], Iter [700/734] Loss: 0.3656
Epoch [5/10], Iter [100/734] Loss: 0.4213
Epoch [5/10], Iter [200/734] Loss: 0.5637
Epoch [5/10], Iter [300/734] Loss: 0.8980
Epoch [5/10], Iter [400/734] Loss: 0.5166
Epoch [5/10], Iter [500/734] Loss: 0.5732
Epoch [5/10], Iter [600/734] Loss: 0.7303
Epoch [5/10], Iter [700/734] Loss: 0.6087
Epoch [6/10], Iter [100/734] Loss: 0.7945
Epoch [6/10], Iter [200/734] Loss: 0.6837
Epoch [6/10], Iter [300/734] Loss: 0.5301
Epoch [6/10], Iter [400/734] Loss: 0.5391
Epoch [6/10], Iter [500/734] Loss: 0.4117
Epoch [6/10], Iter [600/734] Loss: 0.5727
Epoch [6/10], Iter [700/734] Loss: 0.7116
Epoch [7/10], Iter [100/734] Loss: 0.4469
Epoch [7/10], Iter [200/734] Loss: 0.4901
Epoch [7/10], Iter [300/734] Loss: 0.5236
Epoch [7/10], Iter [400/734] Loss: 0.5761
Epoch [7/10], Iter [500/734] Loss: 0.2883
Epoch [7/10], Iter [600/734] Loss: 0.4194
Epoch [7/10], Iter [700/734] Loss: 0.6461
Epoch [8/10], Iter [100/734] Loss: 0.4778
Epoch [8/10], Iter [200/734] Loss: 0.8473
Epoch [8/10], Iter [300/734] Loss: 0.5763
Epoch [8/10], Iter [400/734] Loss: 0.5120

Epoch [8/10], Iter [500/734] Loss: 0.4834
Epoch [8/10], Iter [600/734] Loss: 0.5395
Epoch [8/10], Iter [700/734] Loss: 0.4784
Epoch [9/10], Iter [100/734] Loss: 0.4852
Epoch [9/10], Iter [200/734] Loss: 0.4749
Epoch [9/10], Iter [300/734] Loss: 0.4621
Epoch [9/10], Iter [400/734] Loss: 0.3594
Epoch [9/10], Iter [500/734] Loss: 0.4487
Epoch [9/10], Iter [600/734] Loss: 0.6004
Epoch [9/10], Iter [700/734] Loss: 0.4709
Epoch [10/10], Iter [100/734] Loss: 0.5118
Epoch [10/10], Iter [200/734] Loss: 0.5617
Epoch [10/10], Iter [300/734] Loss: 0.3650
Epoch [10/10], Iter [400/734] Loss: 0.3906
Epoch [10/10], Iter [500/734] Loss: 0.4994
Epoch [10/10], Iter [600/734] Loss: 0.6473
Epoch [10/10], Iter [700/734] Loss: 0.5782
972.239551 seconds
Test Accuracy of the model on the test images: 81 %
Accuracy of bags : 86 %
Accuracy of dresses : 76 %
Accuracy of footwear : 95 %
Accuracy of outerwear : 74 %
Accuracy of skirts : 89 %
Accuracy of tops : 69 %

Summary and conclusions.

When observing the results of our experiments, we found that footwear, bags and skirts consistently performed the best in all of the experiments we conducted out of the categories we chose. In conclusion, the models that we ran performed better than random guess; however, there are ways this model could be improved going forward. One way to do this would be to simply gather more data and possessing more processing power. We could also try tuning the hyper-parameter differently.

Percentage of the code that you found or copied from the internet.

All the codes combined

$$(274 - 46) / (274 + 180) * 100 = 39 \%$$

References.

<https://stackoverflow.com/questions/47850280/fastest-way-to-compute-image-dataset-channel-wise-mean-and-standard-deviation-in>
<http://acberg.com/papers/wheretobuyit2015iccv.pdf>
<https://github.com/flipkart-incubator/fk-visual-search>
https://github.com/Airconaaron/blog_post_visualizing_pytorch_cnn/blob/master/Visualizing%20Learned%20Filters%20in%20PyTorch.ipynb
<https://arxiv.org/pdf/1311.2901.pdf>
<https://stackoverflow.com/questions/47850280/fastest-way-to-compute-image-dataset-channel-wise-mean-and-standard-deviation-in>
https://github.com/amir-jafari/Deep-Learning/blob/master/Pytorch_/6-Conv_Mnist/Conv_Mnist.py
<https://discuss.pytorch.org/t/data-augmentation-in-pytorch/7925>
<https://stackoverflow.com/questions/2301789/read-a-file-in-reverse-order-using-python>
<https://pytorch.org/docs/stable/torchvision/transforms.html>
<https://stackoverflow.com/questions/12984426/python-pil-ioerror-image-file-truncated-with-big-images>
<http://www.apsipa.org/proceedings/2017/CONTENTS/papers2017/14DecThursday/Poster%204/TP-P4.14.pdf>
<https://www.quora.com/How-can-I-calculate-the-size-of-output-of-convolutional-layer>
<https://medium.com/@RaghavPrabhu/understanding-of-convolutional-neural-network-cnn-deep-learning-99760835f148>